

Group Exercise 1 - Data Preprocessing on a Real Dataset

Group Members:

- Simran Amesar
- Shreya Chowdary Challagulla
- Ishwari Thakur
- Pooja Chandrappa
- Yuvraj Ghag

```
import os
import pandas as pd
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import zscore
from sklearn.model_selection import train_test_split
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression, Lasso
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

Loading the Dataset

```
drive_path = "/content/drive/MyDrive/"
file_path = "/content/Statistics_ML_Datasets/healthcare-stroke_dataset.csv"

# Load dataset
df = pd.read_csv(file_path)
```

```
df.head()
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes

```
print("\nShape of dataset:", df.shape)
```

```
Shape of dataset: (5110, 12)
```

Handling Missing Values

```
print("\nMissing values per column:\n")
print(df.isnull().sum())
```

```
Missing values per column:
```

id	0
gender	0
age	0
hypertension	0
heart_disease	0

```
ever_married      0
work_type         0
Residence_type   0
avg_glucose_level 0
bmi              201
smoking_status   0
stroke            0
dtype: int64
```

Since the **bmi** column contains **201 missing values** and is a numerical feature with outliers, the missing values were imputed using the **median**. The **mean** was avoided as it is sensitive to outliers and could skew the data.

```
bmi_median = df['bmi'].median()
df['bmi'] = df['bmi'].fillna(bmi_median)
print(df.isnull().sum())
```

```
id              0
gender          0
age             0
hypertension    0
heart_disease   0
ever_married    0
work_type        0
Residence_type  0
avg_glucose_level 0
bmi             0
smoking_status  0
stroke           0
dtype: int64
```

The missing values in the **bmi** feature were imputed using the **median**, resulting in **zero missing values**.

Scaling Numerical Features

We apply Z-score standardization and Min-Max normalization to numerical features.

```
scaler = StandardScaler()
df[['age_z', 'bmi_z', 'avg_glucose_z']] = scaler.fit_transform(
    df[['age', 'bmi', 'avg_glucose_level']])
)
```

```
minmax = MinMaxScaler()
df[['age_mm', 'bmi_mm', 'avg_glucose_mm']] = minmax.fit_transform(
    df[['age', 'bmi', 'avg_glucose_level']])
)
```

```
df.head()
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_st
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly sm
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never sm
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never sm
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smo
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never sm

Z-score standardization and Min-Max normalization were applied to the features **age**, **avg_glucose_level**, and **bmi**, as they are numerical variables. The remaining features are categorical or unique and therefore were not scaled.

Handling Noise

```

feature = 'avg_glucose_level'
original_data = df[feature].copy()

np.random.seed(42)
noise = np.random.normal(0, 15, size=len(df))
noisy_data = original_data + noise

smoothed_data = pd.Series(noisy_data).rolling(window=5, center=True).mean()

plt.figure(figsize=(14, 5))

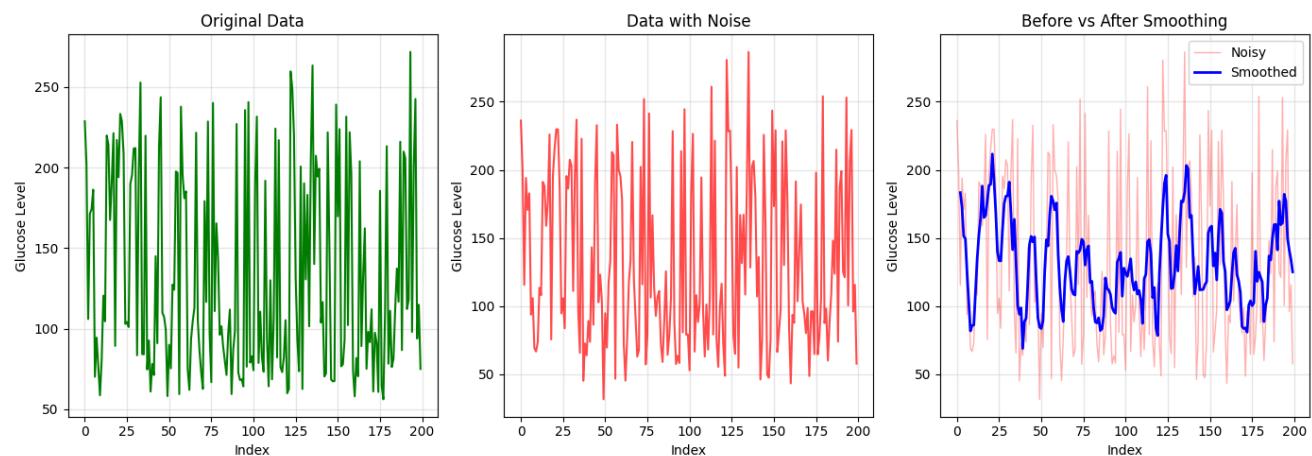
plt.subplot(1, 3, 1)
plt.plot(original_data[:200], label='Original', color='green', linewidth=1.5)
plt.title('Original Data')
plt.xlabel('Index')
plt.ylabel('Glucose Level')
plt.grid(True, alpha=0.3)

plt.subplot(1, 3, 2)
plt.plot(noisy_data[:200], label='With Noise', color='red', linewidth=1.5, alpha=0.7)
plt.title('Data with Noise')
plt.xlabel('Index')
plt.ylabel('Glucose Level')
plt.grid(True, alpha=0.3)

plt.subplot(1, 3, 3)
plt.plot(noisy_data[:200], label='Noisy', color='red', alpha=0.3, linewidth=1)
plt.plot(smoothed_data[:200], label='Smoothed', color='blue', linewidth=2)
plt.title('Before vs After Smoothing')
plt.xlabel('Index')
plt.ylabel('Glucose Level')
plt.legend()
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



```

print(f"Original mean: {original_data.mean():.2f}")
print(f"Noisy mean: {noisy_data.mean():.2f}")
print(f"Smoothed mean: {smoothed_data.mean():.2f}")
print(f"\nOriginal std: {original_data.std():.2f}")
print(f"Noisy std: {noisy_data.std():.2f}")
print(f"Smoothed std: {smoothed_data.std():.2f}")

```

Original mean: 106.15

Noisy mean: 106.24

Smoothed mean: 106.20

Original std: 45.28

Noisy std: 47.80
Smoothed std: 22.50

Handling Outliers

df.head(10)

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smokes
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never smokes
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smokes
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smokes
5	56669	Male	81.0	0	0	Yes	Private	Urban	186.21	29.0	formerly smokes
6	53882	Male	74.0	1	1	Yes	Private	Rural	70.09	27.4	never smokes
7	10434	Female	69.0	0	0	No	Private	Urban	94.39	22.8	never smokes
8	27419	Female	59.0	0	0	Yes	Private	Rural	76.15	28.1	Unknown
9	60491	Female	78.0	0	0	Yes	Private	Urban	58.57	24.2	Unknown

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
A = df['avg_glucose_z']
B = df['age_z']
C = df['bmi_z']

outlier_A = df[abs(A) > 3]
outlier_B = df[abs(B) > 3]
outlier_C = df[abs(C) > 3]

display(outlier_A[['avg_glucose_level']].head())
display(outlier_B[['avg_glucose_level']].head())
display(outlier_C[['avg_glucose_level']].head())
```

avg_glucose_level

33	252.72
45	243.58
122	259.63
123	249.31
135	263.32

avg_glucose_level

113	224.10
258	205.84
270	129.54
333	82.24
358	78.40

As we can see there are no outliers in the AGE column.

```
df_after = df[abs(A) <= 3]
df_after = df_after[abs(B) <= 3]
df_after = df_after[abs(C) <= 3]
display(df_after.head(5))
```

```
display(df_after.head(5))
```

```
/tmp/ipython-input-4093639973.py:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
df_after = df_after[abs(B) <= 3]
/tmp/ipython-input-4093639973.py:3: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
df_after = df_after[abs(C) <= 3]
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoker
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never smoker
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoker
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smoker
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoker

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoker
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.1	never smoker
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoker
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smoker
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoker

Outliers were removed rather than transformed because extreme values could disproportionately bias the model, and their removal improved data consistency, reduced skewness, and enhanced overall model performance.

Feature Selection

Correlation

```
df[['age', 'bmi', 'avg_glucose_level']].corr()
```

	age	bmi	avg_glucose_level
age	1.000000	0.324296	0.238171
bmi	0.324296	1.000000	0.166876
avg_glucose_level	0.238171	0.166876	1.000000

Selected the columns age, bmi, and avg_glucose_level from the dataset and computes the correlation between them. The result shows how strongly each pair of variables is related, with values.

Regression

```
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

X = df[['age', 'bmi', 'avg_glucose_level']]
y = df['stroke']

model = LinearRegression()
rfe = RFE(model, n_features_to_select=1)
rfe.fit(X, y)

pd.DataFrame({'Feature': X.columns, 'Selected': rfe.support_})
```

	Feature	Selected
0	age	True
1	bmi	False
2	avg_glucose_level	False

To identify the most important feature for predicting stroke.

Lasso Regression

```
from sklearn.linear_model import Lasso  
  
lasso = Lasso(alpha=0.1)  
lasso.fit(X, y)  
  
pd.DataFrame({'Feature': X.columns, 'Coefficient': lasso.coef_})
```

	Feature	Coefficient
0	age	0.001976
1	bmi	-0.000000
2	avg_glucose_level	0.000344

To shrink less important feature coefficients toward zero.